

Healthcare Track -- Cortex Code Hands-On Lab

Powered by Cortex Code


About This Lab

In this hands-on lab, you will build a fully functional AI-powered Healthcare Intelligence using **only natural language prompts** in Snowflake's AI-powered development platform, **Cortex Code**. No SQL writing, no YAML editing, no manual UI configuration -- you'll talk to Cortex Code and it will build everything for you.

What you'll build:

- | | |
|-----------------------|---|
| 1. Clinical Database | Patient records, encounters, lab results, and more |
| 2. Cortex Analyst | Ask questions in plain English, get SQL results instantly |
| 3. Cortex Search | Natural language search across clinical PDFs and reports |
| 4. Intelligence Agent | AI orchestrator that routes queries to the right tool automatically |
| 5. Streamlit App | Chat interface, live dashboards, and bed optimization |

Time: 75 minutes hands-on **Prerequisites:** A laptop with a modern browser. No coding experience required. No files to download -- everything comes directly from GitHub.



Step 0: Get Started (8 minutes)

0.1 Create and Log Into Your Snowflake Trial Account

A shared URL will be provided to sign up for a Snowflake trial account.

1. Navigate to the **signup URL** provided by your instructor
2. Ensure **"AI Data Cloud for Enterprise"** is selected (see screenshot below)
3. Fill out the registration form with your details
4. Click **Continue** and complete the signup process

5. Check your email for a confirmation from Snowflake
6. Follow the instructions in the email to activate your account and set your password
7. Log in to your new Snowflake trial account — you should land on the **Snowsight** home page

BUILD FASTER WITH SNOWFLAKE

One platform for your data, apps and AI workflows

AI Data Cloud For EnterpriseCortex Code CLI For Developers

Snowflake AI Data Cloud

Build enterprise-grade AI and apps on your data in minutes

- Gain immediate access to the AI Data Cloud
- Enable your most critical data workloads
- Scale instantly, elastically, and near-infinately across public clouds
- Snowflake is HIPAA, PCI DSS, SOC 1 and SOC 2 Type 2 compliant, and FedRAMP Authorized



Start your 30-day trial with \$400 in free credits

Create a Snowflake account1 / 2

Already have an account? [Sign in](#)

First name

Last name

Work email

Why are you signing up?

By clicking the button below you understand that Snowflake will process your personal information in accordance with its [Privacy Notice](#)

Continue

Countrv: Mexico ▼

0.2 Open Cortex Code

1. Look for the **Cortex Code** icon in the **lower-right corner** of Snowsight
2. Click it -- the Cortex Code panel will open on the right side of the screen
3. You should see a chat input box at the bottom of the panel

What is Cortex Code? Cortex Code is Snowflake's AI-powered development platform. You can ask it to write SQL, create objects, build applications, and more -- all through natural language conversation. Think of it as your AI pair programmer that understands Snowflake natively.

0.3 Connect to the Lab's GitHub Repository

Instead of downloading any files, you'll connect Snowflake directly to the public GitHub repository that contains all the lab materials. Snowflake can read and execute files straight from GitHub -- no ZIP files, no uploads.


Copy and paste this prompt into Cortex Code:

```
Please connect my Snowflake account to a public GitHub repository so I can use it as a file source during this lab.

1. Create a database called HOL_UTILS with schema PUBLIC if it doesn't exist
2. Create an API integration called HOL_GITHUB_API for GitHub HTTPS access, allowing the prefix https://github.com/jfromkamel/
3. Create a Git repository object called SNOWFLAKE_AI_HOL_REPO in the HOL_UTILS.PUBLIC schema pointing to: https://github.com/jfromkamel/snowflake-ai-hol
4. Fetch the latest contents from the repository
5. Confirm everything worked by listing the top-level files and folders available in the main branch
```

What's happening: Cortex Code is creating a live connection between your Snowflake account and the lab's GitHub repository. All scripts, models, and documents will be accessible directly from GitHub -- no manual file handling needed.

Important -- "Always allow" permissions: When Cortex Code tries to run a CREATE statement, you'll see a permission prompt asking whether to allow it. Click the **Allow** dropdown arrow and select **"Allow CREATE in this chat"**. Do the same for any subsequent permission prompts you encounter (e.g., ALTER, INSERT, etc.). This prevents you from having to manually approve every single statement for the rest of the lab.



Expected output: In the Cortex Code response, you should see a folder listing including `healthcare/`, `financial-services/`, `industrial/`, and `README.md`.

Note: This is the only prompt in the lab that requires ACCOUNTADMIN. Cortex Code will handle the role switching automatically.

Fallback SQL for Step 0.3

If Cortex Code can't handle the Git setup, run this SQL directly in a Snowsight worksheet:

```
USE ROLE ACCOUNTADMIN;
CREATE DATABASE IF NOT EXISTS HOL_UTILS;
CREATE SCHEMA IF NOT EXISTS HOL_UTILS.PUBLIC;
CREATE OR REPLACE API INTEGRATION HOL_GITHUB_API
  API_PROVIDER = git_https_api
  API_ALLOWED_PREFIXES = ('https://github.com/jfromkamel/')
  ENABLED = TRUE;
CREATE OR REPLACE GIT REPOSITORY HOL_UTILS.PUBLIC.SNOWFLAKE_AI_HOL_REPO
  API_INTEGRATION = HOL_GITHUB_API
  ORIGIN = 'https://github.com/jfromkamel/snowflake-ai-hol';
ALTER GIT REPOSITORY HOL_UTILS.PUBLIC.SNOWFLAKE_AI_HOL_REPO FETCH;
LS @HOL_UTILS.PUBLIC.SNOWFLAKE_AI_HOL_REPO/branches/main/;
```

0.4 Browse Lab Files in a Workspace

Now that Snowflake is connected to GitHub, create a visual workspace to browse the repo files:

1. In the left sidebar, go to **Projects > Workspaces**
2. Click the **+** next to Workspaces and select **Git workspace**
3. Paste: `https://github.com/jfromkamel/snowflake-ai-hol`
4. Select `HOL_GITHUB_API` as the API Integration
5. Choose **Public repository** for authentication

6. Click **Create**

Expected output: The repo files (`healthcare/`, `financial-services/`, `README.md`) appear in the left-hand file explorer. You can browse any file directly from Snowsight.

Why this matters: Your Snowflake account now has a live, bidirectional connection to GitHub. You can pull the latest changes, create branches, commit edits, push updates back to the remote repo, and resolve merge conflicts -- all directly from Snowsight. This means developers get full Git workflows (branching, pull/push, conflict resolution) natively inside Snowflake without switching tools.

Tip: You can also open `LAB_GUIDE.md` directly in the workspace file explorer (under your industry folder) and follow along from there instead of the PDF -- whichever you prefer.

Step 1: Build the Database and Load Data (10 minutes)

In this step, you'll create a healthcare database with patient records, encounters, diagnoses, lab results, and more -- all by asking Cortex Code to run a single script directly from GitHub.

1.1 Execute the Setup Script

Copy and paste this prompt into Cortex Code:

```
Please execute the setup script from the GitHub repository we connected.
Run this file to create the healthcare database, warehouse, all tables,
internal stages, and load all sample data:
```

```
@HOL_UTILS.PUBLIC.SNOWFLAKE_AI_HOL_REPO/branches/main/healthcare/scripts/setup.sql
```

```
Use EXECUTE IMMEDIATE FROM to run it. After execution, set
HEALTHCARE_HOL_WH as the active warehouse for the rest of our session
and show me a summary table of every table created with its schema and
row count so I can confirm everything loaded correctly.
```

What's happening: Cortex Code is pulling the setup script directly from GitHub and executing it. The script creates: - Database: `HEALTHCARE_HOL_DB` with schemas `CLINICAL` and `OPERATIONS` - Warehouse: `HEALTHCARE_HOL_WH` - 11 tables with realistic patient, provider, and operational data - Internal stages for files you'll use later

Congratulations! You just built an entire database without writing a single line of SQL.

1.2 Explore Your New Database

Take a moment to see what was created. In the left navigation, click **Databases** and expand **HEALTHCARE_HOL_DB > CLINICAL > Tables**. Click on any table (e.g., **BEDS**) to see its details, columns, and a Cortex-generated description that automatically explains what the table contains.

Database Explorer

Why this matters: The Database Explorer gives you instant visibility into every object Cortex Code created -- table schemas, row counts, column types, and AI-generated descriptions -- without writing a single query.

Step 2: Create the Semantic Layer (12 minutes)

Cortex Analyst is Snowflake's text-to-SQL engine. You define your business metrics, table relationships, and business rules once -- then anyone can query your data in plain English without knowing SQL.

In this step, you'll have Cortex Code **set up Cortex Analyst** from scratch.

2.1 Create a Cortex Analyst Semantic View for Clinical Data

This is the "wow" moment. Copy and paste this prompt into Cortex Code:

```
Create a Cortex Analyst semantic view for the clinical data in
HEALTHCARE_HOL_DB.CLINICAL.
```

```
Include these tables: PATIENTS, PROVIDERS, DEPARTMENTS, ENCOUNTERS,
DIAGNOSES, PROCEDURES, MEDICATIONS, LAB_RESULTS, BEDS, and READMISSIONS.
```

```
Business intelligence rules:
```

- Flag "high readmission risk" when `RISK_SCORE > 0.7`
- Flag "extended stay" when `LENGTH_OF_STAY > 7`
- Flag "critical lab" when `RESULT_FLAG = 'Critical'`

```
Make sure it can answer questions like:
```

1. "How many patients are currently admitted?"
2. "What is the average length of stay by department?"
3. "Which patients are at high risk for readmission?"
4. "Show me all critical lab results"

```
Name it PATIENT_CARE and store it in HEALTHCARE_HOL_DB.CLINICAL.
```

```
Register it directly as a semantic view without saving any intermediate
files to a stage.
```

What's happening: Cortex Code is analyzing your table structures, understanding the clinical relationships, and configuring a complete text-to-SQL layer. This would normally take a data engineer hours to write manually.

2.2 Create a Cortex Analyst Semantic View for Claims Analytics

Copy and paste this prompt into Cortex Code:

```
Create a Cortex Analyst semantic view for insurance claims analytics.

Include these tables:
- HEALTHCARE_HOL_DB.OPERATIONS.INSURANCE_CLAIMS
- HEALTHCARE_HOL_DB.CLINICAL.PATIENTS
- HEALTHCARE_HOL_DB.CLINICAL.ENCOUNTERS

Join INSURANCE_CLAIMS to PATIENTS on PATIENT_ID, and to ENCOUNTERS
on ENCOUNTER_ID.

Business intelligence rules:
- Flag "claim denied" when STATUS = 'Denied'
- Calculate approval rate as percentage of claims with STATUS = 'Approved'
- Calculate average reimbursement as APPROVED_AMOUNT / CLAIM_AMOUNT

Make sure it can answer questions like:
1. "What is the total claim amount by insurance type?"
2. "Which claims were denied?"
3. "What is the approval rate by department?"
4. "Show average reimbursement percentage by insurance type"

Name it CLAIMS_ANALYTICS and store it in HEALTHCARE_HOL_DB.CLINICAL.
Register it directly as a semantic view without saving any intermediate
files to a stage.
```

2.3 Explore Your Semantic Model in Snowsight

Step out of Cortex Code for a moment -- let's see what you built in the Snowsight UI.

Navigate to the Analyst:

1. In the left navigation, hover over the **Cortex** icon
2. Click **Analyst**
3. In the **Database** dropdown at the top, select **HEALTHCARE_HOL_DB**, then select the **CLINICAL** schema
4. Find your **PATIENT_CARE** semantic view in the list and click it to open it

Ask a question:

1. In the chat box, type: *"How many patients are currently admitted by department?"*
2. Cortex Analyst generates SQL, runs it, and returns results -- all in one step
3. Click **Add to Verified Queries** below the response to save this as a verified query. Verified queries act as trusted reference answers that improve the model's accuracy over time.

Explore your semantic view:

1. Click the back arrow to return to the semantic view detail page. Take a moment to explore:
2. **Custom Instructions** -- business rules and context you provided (e.g., "high readmission risk = RISK_SCORE > 0.7")
3. **Logical Tables** -- the tables included in this model and their columns
4. **Relationships** -- how tables are joined together
5. **Verified Queries** -- saved question/SQL pairs (including the one you just added) that serve as reference examples
6. **Suggestions** -- the model's recommendations for improving coverage

Note: The Analyst can answer *any* natural language question about the included tables -- not just verified ones. Verified queries simply improve accuracy by giving the model reference examples of correct SQL for common questions.

Ask a follow-up:

1. Now ask: *"What is the average length of stay for cardiology patients?"*
2. Notice the conversation maintains context -- this is a full dialogue, not isolated one-off queries.

Check the monitoring view:

1. Click the **Monitor** tab at the top of the Analyst UI
2. You'll see a history of every question asked, the SQL generated, whether it succeeded, and user feedback. This is how data teams identify gaps in the semantic model and prioritize improvements.

Key takeaway: Anyone in your organization -- a nurse manager, a CMO, an ops analyst -- can now query clinical data in plain English. No SQL, no data engineering ticket, no waiting.

Step 3: Set Up Document Search (8 minutes)

So far, your agent can answer any question about clinical operations, triage protocols, discharge procedures, and readmission policies that lives in structured tables. But a lot of critical knowledge lives somewhere else entirely -- in documents.

When a clinician asks your agent "What is our discharge planning procedure for high-risk patients?", the answer isn't in a database table -- it's in your Clinical Operations Manual. That's the gap Cortex Search fills.

In this step, you'll turn a PDF document into a fully searchable knowledge base. Cortex Search automatically reads the document, breaks it into meaningful sections, and builds a semantic index so the agent can retrieve the right passage for any question -- even when the exact words don't match.

This is what's commonly called **RAG (Retrieval Augmented Generation)** -- but you won't need to configure any of that. Cortex Code handles it in one prompt.

3.1 Bring the PDF into Snowflake

Copy and paste this prompt into Cortex Code:

```
Bring the following PDF from our GitHub repository into
HEALTHCARE_HOL_DB.CLINICAL so we can make it searchable:
```

```
@HOL_UTILS.PUBLIC.SNOWFLAKE_AI_HOL_REPO/branches/main/healthcare/pdfs/Clinical Operations
Manual.pdf
```

3.2 Create the Search Service

Copy and paste this prompt into Cortex Code:

```
Create a Cortex Search service called CLINICAL_DOCS_INFO in
HEALTHCARE_HOL_DB.CLINICAL that makes the clinical operations PDF
searchable with natural language.
```

To do this:

- Temporarily scale up HEALTHCARE_HOL_WH for processing power
- Read and parse the PDF content
- Break it into overlapping chunks for better search accuracy
- Index everything so it can be queried in plain English
- Scale the warehouse back down to SMALL when done

3.3 Preview the Document and Explore the Search Service

First, preview the source PDF:

Before testing search, let's confirm the document loaded correctly and give you a chance to review it. Paste this prompt into Cortex Code:

```
Generate a presigned URL so I can preview the file "Clinical Operations Manual.pdf" in
HEALTHCARE_HOL_DB.CLINICAL.PDF_STAGE -- use an expiry of 3600 seconds.
```

Cortex Code will return a URL. Open it in a new browser tab to read the source document your search service is built on.

Why this matters: Seeing the source document helps you understand why the search returns what it returns. You're not searching a black box -- the content is fully visible and auditable.

Navigate to the Search Playground:

1. In the left navigation, hover over the **Cortex** icon
2. Click **Cortex Search**
3. In the **Database** dropdown, select **HEALTHCARE_HOL_DB**
4. Find **CLINICAL_DOCS_INFO** in the list and click on it to open it
5. You'll land on the service detail page -- click the **Query** tab (or **Playground**) to open the search interface

Run a search: 6. In the search box, type: *"discharge planning procedures"* 7. Review the results -- notice it returns the most relevant passages from the document, not just keyword matches 8. Click into a result to see the full chunk of text it retrieved

Try a semantic search (the real power): 9. Now type: *"triage protocol for chest pain"* 10. Try rephrasing the same concept in different words -- notice the results stay relevant even when your wording doesn't exactly match the document

What you're seeing: This is hybrid search -- combining keyword matching with semantic (meaning-based) matching. It finds the right passage even when someone asks a question differently than the document is written.

Key takeaway: Your agent can now answer policy questions like "What is the readmission prevention protocol for heart failure patients?" -- grounded in your actual documentation, not hallucinated.

Step 4: Build the Intelligence Agent (10 minutes)

Now you'll wire everything together into a single AI agent that automatically routes questions to the right data source.

4.1 Create the Agent

Copy and paste this prompt into Cortex Code:

```
Create a Snowflake Intelligence Agent named HEALTHCARE_AGENT in
HEALTHCARE_HOL_DB.CLINICAL.

Give it three tools:

1. A data analysis tool named PATIENT_CARE_DATA connected to the
PATIENT_CARE semantic model in HEALTHCARE_HOL_DB.CLINICAL.
It answers questions about patients, encounters, diagnoses,
labs, and bed management. Use HEALTHCARE_HOL_WH.

2. A data analysis tool named CLAIMS_DATA connected to the
CLAIMS_ANALYTICS semantic model in HEALTHCARE_HOL_DB.CLINICAL.
It analyzes insurance claims, approval rates, and billing data.
Use HEALTHCARE_HOL_WH.

3. A document search tool named CLINICAL_DOCS connected to the
CLINICAL_DOCS_INFO search service in HEALTHCARE_HOL_DB.CLINICAL.
It searches clinical documentation for protocols and procedures.
```

4.2 Explore the Two UIs

There are two separate places in Snowsight where your agent lives -- and they serve completely different audiences.

View 1: Cortex > Agents (the admin view)

This is where developers and data engineers build and manage agents. Think of it as the control panel.

1. In the left navigation, hover over the **Cortex** icon
2. Click **Agents**
3. Find **HEALTHCARE_AGENT** and click on it

- 4. Open the **Tools** tab -- you'll see all 3 tools listed:
- 5. **PATIENT_CARE_DATA** (Cortex Analyst -- clinical data)
- 6. **CLAIMS_DATA** (Cortex Analyst -- insurance claims)
- 7. **CLINICAL_DOCS** (Cortex Search -- documentation)
- 8. Click on any tool to see its configuration -- the semantic model it points to, the search service it queries, and the warehouse it uses
- 9. Open the **Settings** tab -- here you can edit the agent's system prompt, set its persona, and control what it can and can't do

Who uses this view: Admins and data engineers. This is where you build, update, and govern agents. Business users never need to come here.

View 2: Snowflake Intelligence (the end-user experience)

This is what your business users actually see -- a polished, conversational interface built on top of the same agent.

- 1. In the left navigation, look for **Snowflake Intelligence** (it appears as a standalone section, separate from Cortex)
- 2. If you don't see it, try navigating to **Cortex > Snowflake Intelligence**
- 3. Select **HEALTHCARE_AGENT** from the agent list
- 4. You're now in the end-user chat interface

What's different here compared to Cortex > Agents: - Responses automatically render as **charts or tables** depending on the question type -- no configuration needed - Citations appear below answers, linking back to the exact data source or document passage - Conversations are **threaded** -- users can continue a conversation, branch off, or start fresh - Users can **upload files** (CSV, PDF, XLSX) directly in the chat to add context to their questions - There is an **Extended Thinking** toggle for complex, multi-step analytical questions

Why this matters: The agent you built in Cortex Code is the engine. Snowflake Intelligence is the car. Business users interact with the car -- they never need to know what's under the hood.

4.3 Test the Agent in Snowflake Intelligence

Stay in the Snowflake Intelligence UI. Click **Playground** and try these questions:

Question 1 -- Patient Analysis (routes to clinical data):

Which patients are at high risk for readmission?

Question 2 -- Operational Analysis (complex multi-table query):

For departments with the highest average length of stay, show me the most common diagnoses and the bed occupancy rate.

Question 3 -- Claims Analysis (routes to claims data):

What is the approval rate by insurance type, and which claims were denied?

Question 4 -- Policy Lookup (routes to document search):

What does our documentation say about discharge planning procedures?

Key takeaway: The agent automatically picks the right tool for each question. One interface, every data source, no SQL required.

Step 5: Build a Streamlit App (25 minutes)

You've already built something powerful: an agent accessible through Snowflake Intelligence with conversational AI, auto-generated charts, threaded conversations, and citations. So what does a Streamlit app add on top of that?

Think about a nurse manager who needs a real-time view of bed occupancy and patient risk flags -- in a layout designed for clinical rounds, not open-ended questions. That's a different kind of experience -- a purpose-built application with a fixed layout, specific operational views, and embedded computation like the optimization engine you'll build in this step.

The agent you've created is actually surfaceable in three different ways, and they complement each other:

	Snowflake Intelligence	Streamlit in Snowflake	Cortex Agents REST API
Best for	Open-ended exploration and self-service analytics	Structured operational apps with fixed layouts and custom tools	Embedding the agent anywhere outside Snowflake
Experience	Conversational with auto-charting, threading, and citations	Multi-page app with dashboards, forms, and Python computation	Any interface you build -- chat widgets, portals, integrations

Customization	Agent behavior and instructions	Full control over UI, layout, branding, and logic	Full control -- wire the agent into Slack, Teams, web apps, or mobile
Computation	Answers and analysis from your data	Python on top of live data -- optimization, ML, file processing	Depends on the client you build

The **Cortex Agents REST API** is the same endpoint this Streamlit app will call. That means the agent you built today can also power a Slack bot, a Microsoft Teams integration, an external web portal, or a mobile app -- with no changes to the agent itself.

In this step, you'll build a 3-page app: - **Chat** -- the same agent, now inside your own branded interface - **Dashboard** -- fixed operational views of bed occupancy by department, critical lab alerts, and high-risk patient flags - **Optimization** -- a bed allocation optimizer built in Python that runs directly on your live data

The **big picture**: Snowflake Intelligence is a great out-of-the-box experience. Streamlit is how you build a product around the same intelligence layer -- and the REST API is how you take it everywhere else.

5.1 Create the Chat Page (8 minutes)

Copy and paste this prompt into Cortex Code:

Build a single-page Streamlit app in Snowflake called HEALTHCARE_ASSISTANT_APP in HEALTHCARE_HOL_DB.CLINICAL using HEALTHCARE_HOL_WH.

Use `st.tabs()` to create a horizontal tab bar at the top of the page with three tabs: "Chat", "Dashboard", and "Optimization". All three tabs should be visible and switchable at the top -- no sidebar page navigation.

Start with the Chat tab:

- A sidebar with the title "Healthcare Intelligence" and a brief description of the app
- A chat interface where users can have a conversation with HEALTHCARE_AGENT
- Show the agent's responses in the chat as they arrive
- When the agent returns data, display it as a table below the response
- When the agent returns SQL, show it in a collapsible section
- Keep the conversation history visible throughout the session

Apply a clean clinical design throughout the app:

- Main content background: white (#FFFFFF) with light gray panels (#F1F5F9) and subtle drop shadows on cards -- the surface should feel clean and open
- Sidebar: slate blue (#334155) with white text and a teal left border on the active page -- provides clear navigation contrast without being harsh
- Primary color: soft blue (#2563EB) for headings, buttons, and active states; teal (#0D9488) as a secondary accent
- Typography: sentence case throughout -- no all-caps; use medium weight for section headers and regular weight for body text
- Alerts and critical values: green (#16A34A) for healthy/normal, amber (#D97706) for warnings, red (#DC2626) for critical -- used sparingly
- Data tables: white background, alternating light gray rows (#F8F9FC), status pills (small colored badges) rather than full-row highlights
- KPI cards: white background, rounded corners, thin colored left border indicating status, large bold number in dark slate (#1E293B)
- Charts: white background with the soft blue as the primary series color
- Chat input area: white background matching the main content -- seamlessly part of the page, not a floating element
- Overall feel: a clinical operations dashboard -- precise, trustworthy, and easy to read at a glance during rounds or a hospital staff review

5.2 Open and Test the Chat Page

1. Navigate to the app in Snowsight (Projects > Streamlit > HEALTHCARE_ASSISTANT_APP)

App not loading or showing an error? Copy the full error message, go back to Cortex Code, and paste this prompt:

> My Streamlit app is showing the following error. Please read the app code, identify the cause, and fix it:

[paste your error here] > Cortex Code will read the app files, diagnose the issue, and update the code. Refresh the app once it confirms the fix.

Visual design issues? Cortex Code supports images -- take a screenshot of the problem, go back to Cortex Code, paste the screenshot directly into the chat, and use this prompt:

Here is a screenshot of my Streamlit app showing a visual issue. Please read the app code, identify what is causing it, and fix it.

Cortex Code will read the screenshot, diagnose the layout or styling problem, and update the code. Refresh the app once it confirms the fix.

1. Try asking: "**Which patients are currently admitted to the ICU?**"
2. Watch the agent respond through your app!

5.3 Add the Dashboard Tab (10 minutes)

Copy and paste this prompt into Cortex Code:

Fill in the Dashboard tab of HEALTHCARE_ASSISTANT_APP with the following content.

This page should show live data from HEALTHCARE_HOL_DB.CLINICAL:

TOP ROW - four summary numbers:

- Total Admitted Patients (PATIENTS where STATUS = 'Admitted')
- Bed Occupancy Rate (occupied beds / total beds as percentage)
- Average Length of Stay (from ENCOUNTERS)
- High Readmission Risk Patients (READMISSIONS where RISK_SCORE > 0.7)

MIDDLE ROW - two charts side by side:

- A bar chart of bed occupancy by department, highlighting departments above 80% occupancy
- A pie chart of encounters by type (Inpatient/Outpatient/Emergency)

BOTTOM ROW - a critical alerts table:

- All patients with RISK_SCORE > 0.7, showing: Patient Name, Department, Primary Diagnosis, Risk Score, Days Since Last Discharge, Readmission Reason
- Sort by risk score descending (highest risk first)
- Include a button to download the table as a CSV

Chart styling:

Use plotly for all charts. Color palette: steel blue (#2563EB) as the primary series color, teal (#0D9488) as the secondary. Use green (#16A34A), amber (#D97706), red (#DC2626) for status-based coloring only. All chart backgrounds should be white (FFFFFF) with light gray grid lines. No dark backgrounds on any chart. Position the legend below each chart (orientation="h", y=-0.3) to prevent overlap with chart titles. Add a top margin (t=50) on all charts so the title never crowds the plot area.

5.4 Add the Optimization Tab (7 minutes)

Copy and paste this prompt into Cortex Code:

Fill in the Optimization tab of HEALTHCARE_ASSISTANT_APP with the following content.

Title: "Clinical Intelligence"

Subtitle: "Optimize bed allocation and forecast readmission risk using AI and machine learning on your live clinical data."

At the top of the tab, add a mode selector using radio buttons:

- "Bed Optimizer"
- "Readmission Risk Forecaster"

MODE 1: Bed Optimizer

Show the following controls in a row above the results:

- A multi-select for Department (options from DEPARTMENTS table, plus "All")
- A slider for Occupancy Alert Threshold: 60% to 95% (default 80%)
- A toggle: "Include patients nearing discharge (LOS > 5 days)"
- A steel blue "Run Optimizer" button

When clicked:

1. Pulls current bed status from BEDS and department occupancy from ENCOUNTERS
2. Finds departments above the selected occupancy threshold
3. If the discharge toggle is on, flags patients with LENGTH_OF_STAY > 5 as near-discharge candidates
4. Uses linear programming (scipy) to suggest moves that balance occupancy while keeping patients in the correct bed type (ICU stays ICU, Pediatric stays Pediatric, etc.)

Display results as:

- Three KPI cards: Departments Over Threshold, Beds That Can Be Freed, Patients Flagged for Discharge Review
- A plotly grouped bar chart showing before/after occupancy by department, with a horizontal reference line at the selected threshold
- A recommended actions table: Patient ID, From Department, To Department, Bed Type, Reason -- sorted by urgency
- Apply the app styling: white background, clinical blue/teal palette, plotly white chart backgrounds

MODE 2: Readmission Risk Forecaster

Show the following controls:

- A dropdown for Department (from DEPARTMENTS table)
- Radio buttons for Forecast Horizon: 7 days / 14 days / 30 days
- A steel blue "Generate Forecast" button

When clicked:

1. Pull the daily count of high-risk patients (RISK_SCORE > 0.7) from READMISSIONS joined with ENCOUNTERS, grouped by department and date
2. Use SNOWFLAKE.ML.FORECAST to predict the number of high-risk patients expected in that department over the selected horizon
3. Retrieve forecast values with upper and lower confidence bounds

Display results as:

- A headline KPI card: "Patients at High Readmission Risk (Forecasted)" for the selected horizon period
- A plotly line chart with:
 - Historical high-risk patient count (solid blue line, #2563EB)
 - Forecasted count (dotted line with light blue confidence band)
 - X-axis: dates, Y-axis: patient count
 - Legend below the chart, white background
- A risk summary table: Date, Forecasted High-Risk Count, vs Current Average, Change (% above/below baseline)
- A small note: "Forecast based on [N] days of historical data. Confidence intervals widen with shorter history."

Apply the same styling as the rest of the app throughout: white background, soft blue primary (#2563EB), teal accent (#0D9488), plotly white chart backgrounds, sentence case, no all-caps.

Add scipy to the app's dependencies.

5.5 Test Your Complete App

1. **Chat** -- Ask "What are the most common diagnoses this month?"
2. **Dashboard** -- Review occupancy rates and readmission risk alerts
3. **Optimization** -- Click "Run Optimization" and see bed rebalancing recommendations

Congratulations! You've built a complete AI-powered healthcare application -- clinical database, semantic layer, document search, intelligent agent, and a polished UI -- all by talking to Cortex Code.